

```
In [1]: 1 import pandas as pd
```

Ans1) Q1. Import the dataset and find the basic descriptive statistics by using many functions and develop insights (5)

```
In [2]: 1 df=pd.read_excel("E:\TRISEMESTER 4\RPFM final.xlsx")
```

```
In [3]: 1 df
```

Out[3]:

	make	segment	model	fuel_type	...	is_brake_assist	is_power_door_locks	is_central_locking	is_power_steering	is_driver_seat_
0	1	A	M1	CNG	...	No	No	No	Yes	
3	1	A	M1	CNG	...	No	No	No	Yes	
3	1	A	M1	CNG	...	No	No	No	Yes	
2	1	C1	M2	Petrol	...	Yes	Yes	Yes	Yes	
3	2	A	M3	Petrol	...	No	Yes	Yes	Yes	
.	
4	2	A	M3	Petrol	...	No	Yes	Yes	Yes	
3	1	A	M1	CNG	...	No	No	No	Yes	
3	1	A	M1	CNG	...	No	No	No	Yes	
4	1	B2	M6	Petrol	...	Yes	Yes	Yes	Yes	
4	3	C2	M4	Diesel	...	Yes	Yes	Yes	Yes	

In [4]: 1 print(df)

	policy_id	policy_tenure	age_of_car	age_of_policyholder	area_cluster	\
0	ID00001	0.515874	0.05	0.644231	C1	
1	ID00002	0.672619	0.02	0.375000	C2	
2	ID00003	0.841110	0.02	0.384615	C3	
3	ID00004	0.900277	0.11	0.432692	C4	
4	ID00005	0.596403	0.11	0.634615	C5	
...	
58587	ID58588	0.355089	0.13	0.644231	C8	
58588	ID58589	1.199642	0.02	0.519231	C14	
58589	ID58590	1.162273	0.05	0.451923	C5	
58590	ID58591	1.236307	0.14	0.557692	C8	
58591	ID58592	0.124429	0.02	0.442308	C8	

	population_density	make	segment	model	fuel_type	... is_brake_assist	\
0	4990	1	A	M1	CNG	...	No
1	27003	1	A	M1	CNG	...	No
2	4076	1	A	M1	CNG	...	No
3	21622	1	C1	M2	Petrol	...	Yes
4	34738	2	A	M3	Petrol	...	No
...
58587	8794	2	A	M3	Petrol	...	No
58588	7788	1	A	M1	CNG	...	No
58589	34738	1	A	M1	CNG	...	No
58590	8794	1	B2	M6	Petrol	...	Yes
58591	8794	3	C2	M4	Diesel	...	Yes

	is_power_door_locks	is_central_locking	is_power_steering	\
0	No	No	Yes	
1	No	No	Yes	
2	No	No	Yes	
3	Yes	Yes	Yes	
4	Yes	Yes	Yes	
...	
58587	Yes	Yes	Yes	
58588	No	No	Yes	
58589	No	No	Yes	
58590	Yes	Yes	Yes	
58591	Yes	Yes	Yes	

	is_driver_seat_height_adjustable	is_day_night_rear_view_mirror	is_ecw	\
0		No	No	No
1		No	No	No
2		No	No	No
3		Yes	Yes	Yes
4		No	Yes	Yes
...	
58587		No	Yes	Yes
58588		No	No	No
58589		No	No	No
58590		Yes	Yes	Yes
58591		Yes	No	Yes

	is_speed_alert	ncap_rating	is_claim
0	Yes	0	0
1	Yes	0	0
2	Yes	0	0
3	Yes	2	0
4	Yes	2	0
...
58587	Yes	2	0
58588	Yes	0	0
58589	Yes	0	0
58590	Yes	2	0
58591	Yes	3	0

[58592 rows x 44 columns]

```
In [5]: 1 df.head() # These are the first 5 rows of the data.
```

Out[5]:

	tenure	age_of_car	age_of_policyholder	area_cluster	population_density	make	segment	model	fuel_type	...	is_brake_assist	is_power_steering
5874		0.05	0.644231	C1	4990	1	A	M1	CNG	...	No	No
2619		0.02	0.375000	C2	27003	1	A	M1	CNG	...	No	No
1110		0.02	0.384615	C3	4076	1	A	M1	CNG	...	No	No
3277		0.11	0.432692	C4	21622	1	C1	M2	Petrol	...	Yes	Yes
3403		0.11	0.634615	C5	34738	2	A	M3	Petrol	...	No	No

```
In [6]: 1 df.tail() # These are the last 5 rows of the data.
```

Out[6]:

	tenure	age_of_car	age_of_policyholder	area_cluster	population_density	make	segment	model	fuel_type	...	is_brake_assist	is_power_steering
89		0.13	0.644231	C8	8794	2	A	M3	Petrol	...	No	No
42		0.02	0.519231	C14	7788	1	A	M1	CNG	...	No	No
73		0.05	0.451923	C5	34738	1	A	M1	CNG	...	No	No
07		0.14	0.557692	C8	8794	1	B2	M6	Petrol	...	Yes	Yes
29		0.02	0.442308	C8	8794	3	C2	M4	Diesel	...	Yes	Yes

```
In [8]: 1 df.shape # In the data we have total 58592 number of observations in rows and 44 columns.
```

Out[8]: (58592, 44)

```
In [9]: 1 df.columns
```

Out[9]: Index(['policy_id', 'policy_tenure', 'age_of_car', 'age_of_policyholder', 'area_cluster', 'population_density', 'make', 'segment', 'model', 'fuel_type', 'max_torque', 'max_power', 'engine_type', 'airbags', 'is_esc', 'is_adjustable_steering', 'is_tpms', 'is_parking_sensors', 'is_parking_camera', 'rear_brakes_type', 'displacement', 'cylinder', 'transmission_type', 'gear_box', 'steering_type', 'turning_radius', 'length', 'width', 'height', 'gross_weight', 'is_front_fog_lights', 'is_rear_window_wiper', 'is_rear_window_washer', 'is_rear_window_defogger', 'is_brake_assist', 'is_power_door_locks', 'is_central_locking', 'is_power_steering', 'is_driver_seat_height_adjustable', 'is_day_night_rear_view_mirror', 'is_ecw', 'is_speed_alert', 'ncap_rating', 'is_claim'], dtype='object')

These are the names of all columns which we have in our dataset.

In [10]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 58592 entries, 0 to 58591
Data columns (total 44 columns):
#   Column                                  Non-Null Count  Dtype
---  -
0   policy_id                             58592 non-null  object
1   policy_tenure                         58592 non-null  float64
2   age_of_car                           58592 non-null  float64
3   age_of_policyholder                  58592 non-null  float64
4   area_cluster                         58592 non-null  object
5   population_density                   58592 non-null  int64
6   make                                 58592 non-null  int64
7   segment                             58592 non-null  object
8   model                               58592 non-null  object
9   fuel_type                           58592 non-null  object
10  max_torque                           58592 non-null  object
11  max_power                            58592 non-null  object
12  engine_type                         58592 non-null  object
13  airbags                             58592 non-null  int64
14  is_esc                              58592 non-null  object
15  is_adjustable_steering               58592 non-null  object
16  is_tpms                             58592 non-null  object
17  is_parking_sensors                   58592 non-null  object
18  is_parking_camera                    58592 non-null  object
19  rear_brakes_type                     58592 non-null  object
20  displacement                         58592 non-null  int64
21  cylinder                             58592 non-null  int64
22  transmission_type                    58592 non-null  object
23  gear_box                             58592 non-null  int64
24  steering_type                        58592 non-null  object
25  turning_radius                       58592 non-null  float64
26  length                               58592 non-null  int64
27  width                                58592 non-null  int64
28  height                               58592 non-null  int64
29  gross_weight                         58592 non-null  int64
30  is_front_fog_lights                  58592 non-null  object
31  is_rear_window_wiper                 58592 non-null  object
32  is_rear_window_washer                58592 non-null  object
33  is_rear_window_defogger              58592 non-null  object
34  is_brake_assist                      58592 non-null  object
35  is_power_door_locks                  58592 non-null  object
36  is_central_locking                   58592 non-null  object
37  is_power_steering                    58592 non-null  object
38  is_driver_seat_height_adjustable     58592 non-null  object
39  is_day_night_rear_view_mirror        58592 non-null  object
40  is_ecw                               58592 non-null  object
41  is_speed_alert                       58592 non-null  object
42  ncap_rating                          58592 non-null  int64
43  is_claim                             58592 non-null  int64
dtypes: float64(4), int64(12), object(28)
memory usage: 19.7+ MB
```

This is the basic info of the data which is showing which column is having which kind of data.

Int64 = complete number. Example= 4,5 float64 = Number in points which is not complete. For example= 4.5, 5.2 Object = which is text+number

```
In [11]: 1 # For Checking missing value
          2 print(df.isnull().sum())
```

```

policy_id 0
policy_tenure 0
age_of_car 0
age_of_policyholder 0
area_cluster 0
population_density 0
make 0
segment 0
model 0
fuel_type 0
max_torque 0
max_power 0
engine_type 0
airbags 0
is_esc 0
is_adjustable_steering 0
is_tpms 0
is_parking_sensors 0
is_parking_camera 0
rear_brakes_type 0
displacement 0
cylinder 0
transmission_type 0
gear_box 0
steering_type 0
turning_radius 0
length 0
width 0
height 0
gross_weight 0
is_front_fog_lights 0
is_rear_window_wiper 0
is_rear_window_washer 0
is_rear_window_defogger 0
is_brake_assist 0
is_power_door_locks 0
is_central_locking 0
is_power_steering 0
is_driver_seat_height_adjustable 0
is_day_night_rear_view_mirror 0
is_ecw 0
is_speed_alert 0
ncap_rating 0
is_claim 0
dtype: int64

```

This is showing that we don't have any null value in our data set.

```
In [12]: 1 df.describe()
```

Out[12]:

in_density	make	airbags	displacement	cylinder	gear_box	turning_radius	length	width	
92.000000	58592.000000	58592.000000	58592.000000	58592.000000	58592.000000	58592.000000	58592.000000	58592.000000	58592.000000
26.858667	1.763722	3.137066	1162.355851	3.626963	5.245443	4.852893	3850.476891	1672.233667	155
60.174792	1.136988	1.832641	266.304786	0.483616	0.430353	0.228061	311.457119	112.089135	7
90.000000	1.000000	1.000000	796.000000	3.000000	5.000000	4.500000	3445.000000	1475.000000	147
12.000000	1.000000	2.000000	796.000000	3.000000	5.000000	4.600000	3445.000000	1515.000000	147
94.000000	1.000000	2.000000	1197.000000	4.000000	5.000000	4.800000	3845.000000	1735.000000	153
03.000000	3.000000	6.000000	1493.000000	4.000000	5.000000	5.000000	3995.000000	1755.000000	163
30.000000	5.000000	6.000000	1498.000000	4.000000	6.000000	5.200000	4300.000000	1811.000000	182



In [13]: 1 print(df.describe())

```
      policy_tenure  age_of_car  age_of_policyholder  population_density \
count  58592.000000  58592.000000  58592.000000  58592.000000
mean      0.611246      0.069424      0.469420      18826.858667
std       0.414156      0.056721      0.122886      17660.174792
min       0.002735      0.000000      0.288462       290.000000
25%      0.210250      0.020000      0.365385      6112.000000
50%      0.573792      0.060000      0.451923      8794.000000
75%      1.039104      0.110000      0.548077     27003.000000
max       1.396641      1.000000      1.000000     73430.000000

      make  airbags  displacement  cylinder  gear_box \
count  58592.000000  58592.000000  58592.000000  58592.000000  58592.000000
mean      1.763722      3.137066     1162.355851      3.626963      5.245443
std       1.136988      1.832641      266.304786      0.483616      0.430353
min       1.000000      1.000000      796.000000      3.000000      5.000000
25%      1.000000      2.000000      796.000000      3.000000      5.000000
50%      1.000000      2.000000      1197.000000      4.000000      5.000000
75%      3.000000      6.000000     1493.000000      4.000000      5.000000
max       5.000000      6.000000     1498.000000      4.000000      6.000000

      turning_radius  length  width  height  gross_weight \
count  58592.000000  58592.000000  58592.000000  58592.000000  58592.000000
mean      4.852893     3850.476891     1672.233667     1553.33537      1385.276813
std       0.228061      311.457119      112.089135       79.62227      212.423085
min       4.500000     3445.000000     1475.000000     1475.00000      1051.000000
25%      4.600000     3445.000000     1515.000000     1475.00000      1185.000000
50%      4.800000     3845.000000     1735.000000     1530.00000      1335.000000
75%      5.000000     3995.000000     1755.000000     1635.00000      1510.000000
max       5.200000     4300.000000     1811.000000     1825.00000      1720.000000

      ncap_rating  is_claim
count  58592.000000  58592.000000
mean      1.759950      0.063968
std       1.389576      0.244698
min       0.000000      0.000000
25%      0.000000      0.000000
50%      2.000000      0.000000
75%      3.000000      0.000000
max       5.000000      1.000000
```

```
1 This is showing the basic descriptive statistics of the data.
2 count is for total number of observation.
3 mean is for mean(average) of data.
4 std is for standard deviation which is showing that how much variance in the data from the mean value.
5 min is for smallest value.
6 max is for the largest value.
7 25% is 1 quartile.
8 50% is half of the data.
9 75% is 1/3 of data.
10
```

1 **# Ans3. Develop predictive model using any classification machine learning technique then validate the model by splitting the data into training and testing (15)**

In [26]: 1 #Library- Lable encoder
2 from sklearn.preprocessing import LabelEncoder
3 le_is_front_fog_lights=LabelEncoder()
4 le_is_central_locking=LabelEncoder()
5
6 input=df.drop(['is_claim'],axis='columns')
7

In [27]: 1 input.columns

Out[27]: Index(['policy_id', 'policy_tenure', 'age_of_car', 'age_of_policyholder', 'area_cluster', 'population_density', 'make', 'segment', 'model', 'fuel_type', 'max_torque', 'max_power', 'engine_type', 'airbags', 'is_esc', 'is_adjustable_steering', 'is_tpms', 'is_parking_sensors', 'is_parking_camera', 'rear_brakes_type', 'displacement', 'cylinder', 'transmission_type', 'gear_box', 'steering_type', 'turning_radius', 'length', 'width', 'height', 'gross_weight', 'is_front_fog_lights', 'is_rear_window_wiper', 'is_rear_window_washer', 'is_rear_window_defogger', 'is_brake_assist', 'is_power_door_locks', 'is_central_locking', 'is_power_steering', 'is_driver_seat_height_adjustable', 'is_day_night_rear_view_mirror', 'is_ecw', 'is_speed_alert', 'ncap_rating'], dtype='object')

In [28]: 1 output=df.is_claim

In [51]: 1 input['is_front_fog_lights_n'] = le_Department.fit_transform(input['is_front_fog_lights'])
2 input['is_central_locking_n'] = le_salary.fit_transform(input['is_central_locking'])

In [67]: 1 input

Out[67]:

s	is_central_locking	is_power_steering	is_driver_seat_height_adjustable	is_day_night_rear_view_mirror	is_ecw	is_speed_alert	ncap
o	No	Yes	No	No	No	Yes	
o	No	Yes	No	No	No	Yes	
o	No	Yes	No	No	No	Yes	
s	Yes	Yes	Yes	Yes	Yes	Yes	
s	Yes	Yes	No	Yes	Yes	Yes	
..	
s	Yes	Yes	No	Yes	Yes	Yes	
o	No	Yes	No	No	No	Yes	
o	No	Yes	No	No	No	Yes	
s	Yes	Yes	Yes	Yes	Yes	Yes	
s	Yes	Yes	Yes	No	Yes	Yes	

In [68]: 1 input_n=input.drop(['is_front_fog_lights', 'is_central_locking'],axis='columns')
2 input_n.columns
3

Out[68]: Index(['policy_id', 'policy_tenure', 'age_of_car', 'age_of_policyholder', 'area_cluster', 'population_density', 'make', 'segment', 'model', 'fuel_type', 'max_torque', 'max_power', 'engine_type', 'airbags', 'is_esc', 'is_adjustable_steering', 'is_tpms', 'is_parking_sensors', 'is_parking_camera', 'rear_brakes_type', 'displacement', 'cylinder', 'transmission_type', 'gear_box', 'steering_type', 'turning_radius', 'length', 'width', 'height', 'gross_weight', 'is_rear_window_wiper', 'is_rear_window_washer', 'is_rear_window_defogger', 'is_brake_assist', 'is_power_door_locks', 'is_power_steering', 'is_driver_seat_height_adjustable', 'is_day_night_rear_view_mirror', 'is_ecw', 'is_speed_alert', 'ncap_rating', 'is_front_fog_lights_n', 'is_central_locking_n'], dtype='object')

```
In [69]: 1 input_n_ =input.drop([ 'policy_id', 'policy_tenure', 'age_of_car', 'age_of_policyholder',
2      'area_cluster', 'population_density', 'make', 'segment', 'model',
3      'fuel_type', 'max_torque', 'max_power', 'engine_type',
4      'is_esc', 'is_adjustable_steering', 'is_tpms', 'is_parking_sensors',
5      'is_parking_camera', 'rear_brakes_type', 'displacement',
6      'transmission_type', 'gear_box', 'steering_type', 'turning_radius',
7      'length', 'width', 'height', 'gross_weight', 'is_rear_window_wiper',
8      'is_rear_window_washer', 'is_rear_window_defogger', 'is_brake_assist',
9      'is_power_door_locks', 'is_power_steering',
10     'is_driver_seat_height_adjustable', 'is_day_night_rear_view_mirror',
11     'is_ecw', 'is_speed_alert'],axis='columns')
12 input_n_.columns
13
```

```
Out[69]: Index(['airbags', 'cylinder', 'is_front_fog_lights', 'is_central_locking',
      'ncap_rating', 'is_front_fog_lights_n', 'is_central_locking_n'],
      dtype='object')
```

```
In [70]: 1 Input =input_n_.drop(['is_front_fog_lights','is_central_locking' ],axis='columns')
```

```
In [72]: 1 Input
```

Out[72]:

	airbags	cylinder	ncap_rating	is_front_fog_lights_n	is_central_locking_n
0	2	3	0	0	0
1	2	3	0	0	0
2	2	3	0	0	0
3	2	4	2	1	1
4	2	3	2	0	1
...
58587	2	3	2	0	1
58588	2	3	0	0	0
58589	2	3	0	0	0
58590	2	4	2	1	1
58591	6	4	3	1	1

58592 rows × 5 columns

```
In [58]: 1 X=Input
2 y=output
```

```
In [59]: 1 #Train test split
2 from sklearn.model_selection import train_test_split
3 X_train, X_test, y_train, y_test = train_test_split(X,y,train_size=0.7)
```

Here we have take the input range which is X have variable under it ('airbags', 'cylinder', 'ncap_rating', 'is_front_fog_lights_n', 'is_central_locking_n')

For output range we have taken (is_claim)

After this i have divie the model in 70% and 30%. In this i will give 70% of the data to the model and model will predict rest of 30% data. We will train th model by 70% and take the result of rest 30% of the data. "#70% use to build the model and rest 30% is to test/ validate the modelX_train, X_test, y_train, y_test = train_test_split(X,y,train_size=0.7) #70% use to build the model and rest 30% is to test/ validate the model"

```
In [60]: 1 from sklearn.linear_model import LogisticRegression
2 model= LogisticRegression()
```

```
In [61]: 1 model.fit(X_train,y_train)
```

```
Out[61]: LogisticRegression
LogisticRegression()
```



```
In [78]: 1 model.predict(X_test)
```

```
Out[78]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [84]: 1 model.score(X_test,y_test)
```

```
Out[84]: 0.9376493343952668
```

1 # Ans4) Find the accuracy rate of the model (5)

```
In [85]: 1 from sklearn.metrics import accuracy_score,classification_report,confusion_matrix,roc_curve, auc
```

```
In [86]: 1 #Evaluate the model
2 y_pred = model.predict(X_test)
3 accuracy = accuracy_score(y_test, y_pred)
4 print("Accuracy:{:.2f}%".format(accuracy * 100))
```

```
Accuracy:93.76%
```

Here the accuracy score of my model is 93.76% which mean tha our model is able to predict the data with very high accuracy,the chances of wrong prediction is less because the accuracy is very high.

1 # Ans5) Explain the confusion matrix the model

```
In [89]: 1 # evaluate the model
2 print("confusion Matrix:\n", confusion_matrix(y_test, y_pred))
3 print("\nclacification Report: \n", classification_report(y_test, y_pred))
```

```
confusion Matrix:
```

```
[[16482   0]
 [ 1096   0]]
```

```
classification Report:
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	16482
1	0.00	0.00	0.00	1096
accuracy			0.94	17578
macro avg	0.47	0.50	0.48	17578
weighted avg	0.88	0.94	0.91	17578

C:\Users\Gaurav Singh Rawat\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Gaurav Singh Rawat\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Gaurav Singh Rawat\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

Here we have the confusion metrics which showing that model is able to predict: 16482 as true negative which mean that 16482 are correct predict by model which doesn't raise claim. 1096 as false negative wich mean that 1096 are incorrecly prect by model which has raise the claim but modek predict that they didn't raise the claim. 0 as false positive which mean model have predict 0 as those who have not raise the claim and correctly predict as per the data. 0 as true positive which mean model have predict 0 as those who raise the claim and correctly predict as per the data.

Precision= It is showing the positively predict data with total number of positive data. Recall= It show that positively predict data with total number of correctly data whether it is true negative or true positive.

1 # Ans 2) Import both the libraries matplotlib and seaborn to plot the cross tabulations and heatmap of the categorical and ration data and mention your insights

```
In [90]: 1 import seaborn as sns

In [91]: 1 import matplotlib as mtlyb


In [92]: 1 pd.crosstab(df.segment,df.is_claim)
```

Out[92]:

is_claim	0	1
segment		
A	16275	1046
B1	3929	244
B2	17058	1256
C1	3329	228
C2	13117	901
Utility	1136	73

```
In [93]: 1 pd.crosstab(df.segment,df.is_claim).plot(kind='bar')
```

Out[93]: <Axes: xlabel='segment'>



It is showing the cross tabulation of different segment with response to in which segment how many have raise the claim 1 is for claim 0 is for Not raise the claim In segment Utility minimum of claim has been raised.Maximum claim has been raised in B2.


```
In [97]: 1 pd.crosstab(df.model,df.is_claim)
```

Out[97]:

is_claim	0	1
model		
M1	14030	918
M10	1136	73
M11	348	15
M2	1000	80
M3	2245	128
M4	13117	901
M5	1482	116
M6	12837	939
M7	2739	201
M8	3929	244
M9	1981	133

```
In [98]: 1 pd.crosstab(df.model,df.is_claim).plot(kind='bar')
```


Out[98]: <Axes: xlabel='model'>



```
1 Here it is showing the claim response on the basis of model in this max claim is by M1 and min is in M11.
```

```
In [100]: 1 pd.crosstab(df.fuel_type,df.is_claim)
          2 pd.crosstab(df.fuel_type,df.is_claim).plot(kind='bar')
```

Out[100]: <Axes: xlabel='fuel_type'>



1 Here it is showing the claim response on the basis of Fuel type in this max claim is of petrol and min is in Diesel.

```
In [107]: 1 # Compute correlation matrix
          2 corr = df[['age_of_car', 'age_of_policyholder', 'population_density', 'is_claim']].corr()
```

```
In [108]: 1 # Heatmap of the correlation matrix
          2 plt.figure(figsize=(10, 8))
          3 sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f')
          4 plt.title('Correlation Matrix')
          5 plt.show()
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[108], line 2
      1 # Heatmap of the correlation matrix
----> 2 plt.figure(figsize=(10, 8))
      3 sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f')
      4 plt.title('Correlation Matrix')

NameError: name 'plt' is not defined
```

```
In [ ]: 1
```